

Feedback Propagation in a Solver-Critic LLM System

What internal representations change when an LLM solver accepts or rejects critic feedback, and are those representations causally involved in feedback uptake?

