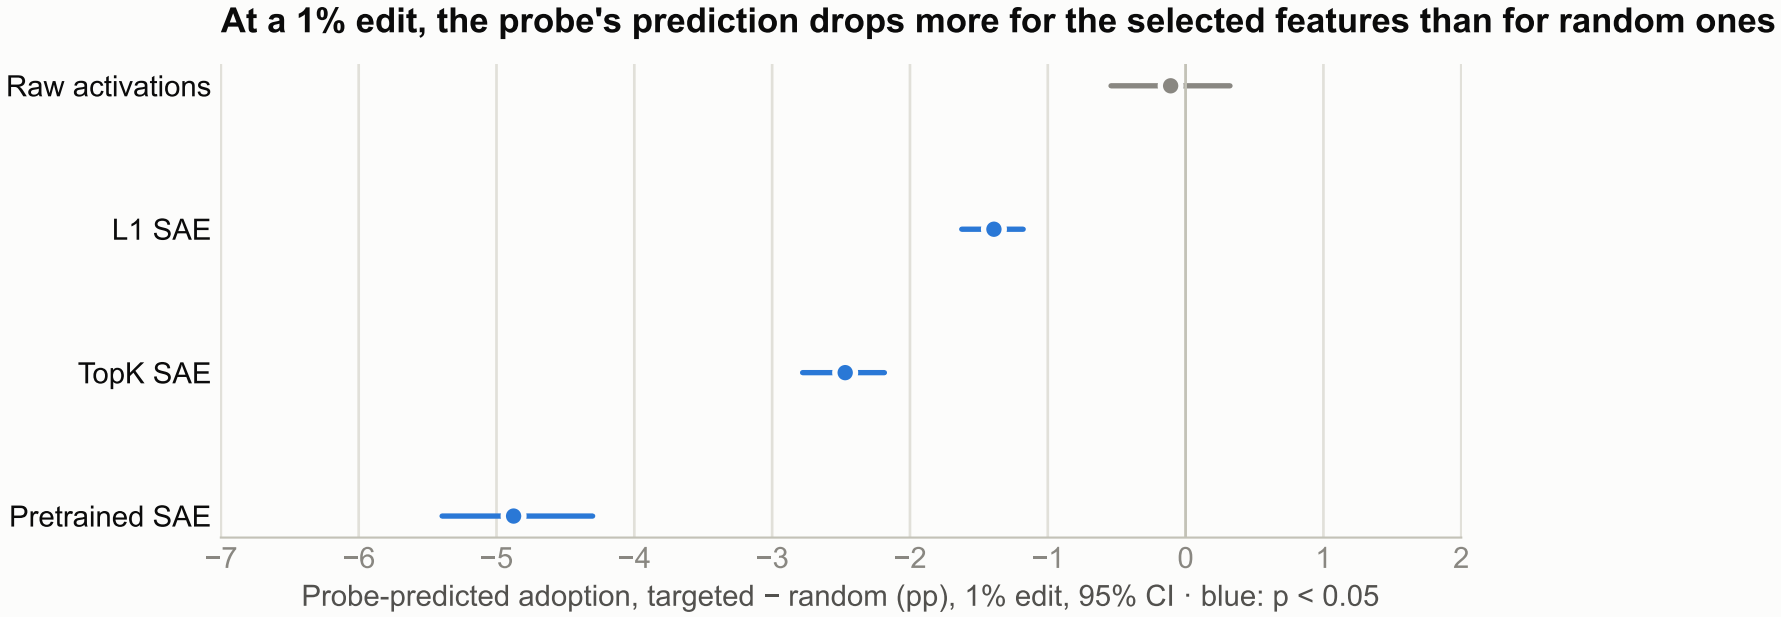
At a 1% edit, the probe's prediction drops more for the selected features than for random ones
Raw activations
L1 SAE
TopK SAE
Pretrained SAE
−7
−6
−5
−4
−3
−2
−1
0
1
2
Probe-predicted adoption, targeted − random (pp), 1% edit, 95% CI · blue: p < 0.05