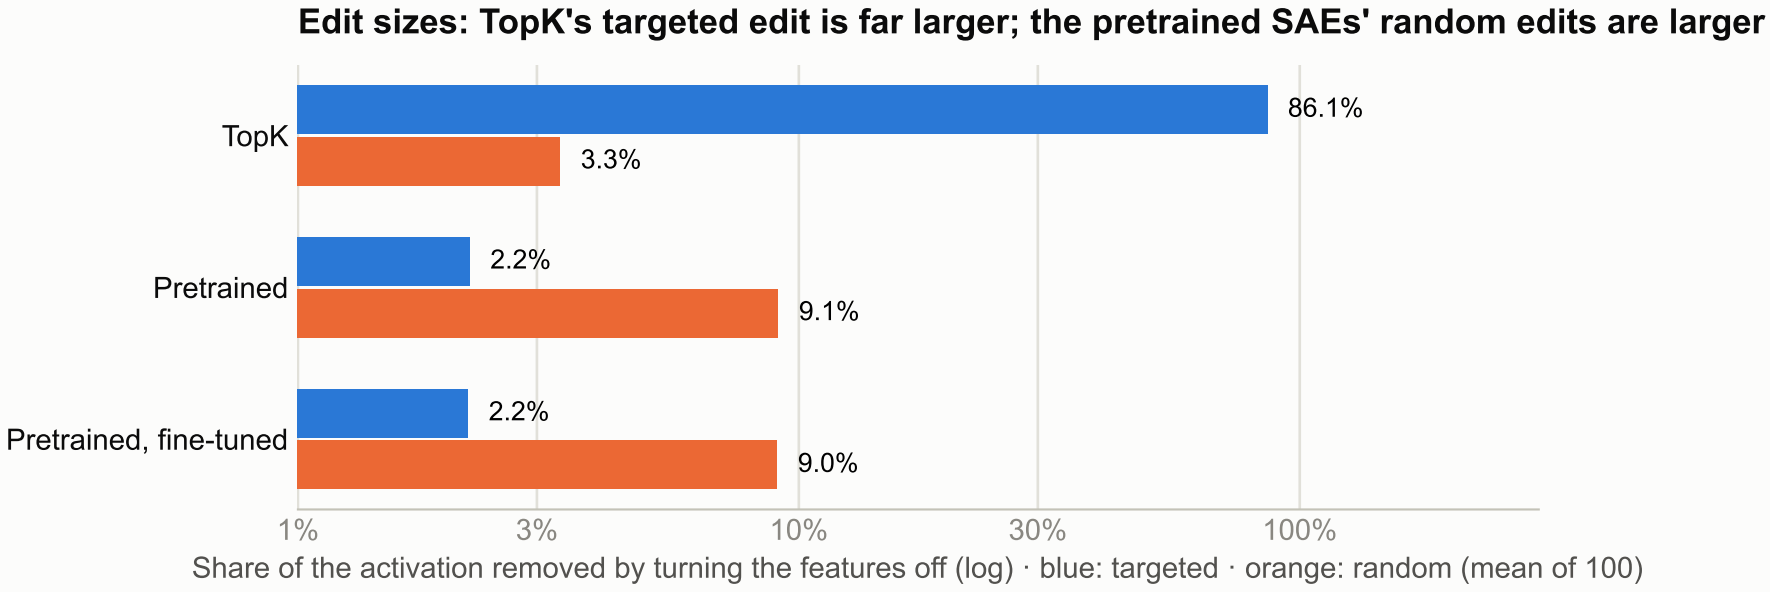
Edit sizes: TopK's targeted edit is far larger; the pretrained SAEs' random edits are larger
TopK
Pretrained
Pretrained, fine-tuned
86.1%
3.3%
2.2%
9.1%
2.2%
9.0%
1%
3%
10%
30%
100%
Share of the activation removed by turning the features off (log) · blue: targeted · orange: random (mean of 100)